

# GitOps

# "CIOps"

Let's say someone uses a hosted CI product like Bitbucket Pipelines and they want to deploy applications from CI to Kubernetes

# "CIOps" the Anti-Pattern

- Give the CI access to the Kubernetes API (security risks)
- CI needs to be configured to the right cluster
- CI needs to have up-to-date credentials
- New cluster => re-configuring multiple pipelines



```
kubectl set image
```

```
helm upgrade
```

# But what about In-Cluster "CI Ops"

- Avoids some security issues, since the CI is inside the cluster
- The configuration problem goes away

# Other issues

- Need to have enough resources to run your builds
- Manage build history
- With multiple clusters images gets independently built per cluster. Not 100% identical image builds, and ruins the idea of container images. Builds needs to be truly reproducible.

# Deployments failes

## How to find correct version to rollback to?

- Trace your build logs
- Kubernetes workload revisions

## What if the production cluster fails?

- How to tell what version of each app you need to deploy?
- Need to re-run all of your CI jobs for that cluster
- How to ensure the CI job will deploy to the right cluster?

# GitOps in 1 slide

System development/managment pattern:

- Git as the **SINGLE** source of truth of a system
- Git as the **SINGLE** place where we operate (create, change and destroy) **ALL** environments
- **ALL changes are obervable/verifiable**

# GitOps - Operations done through Git

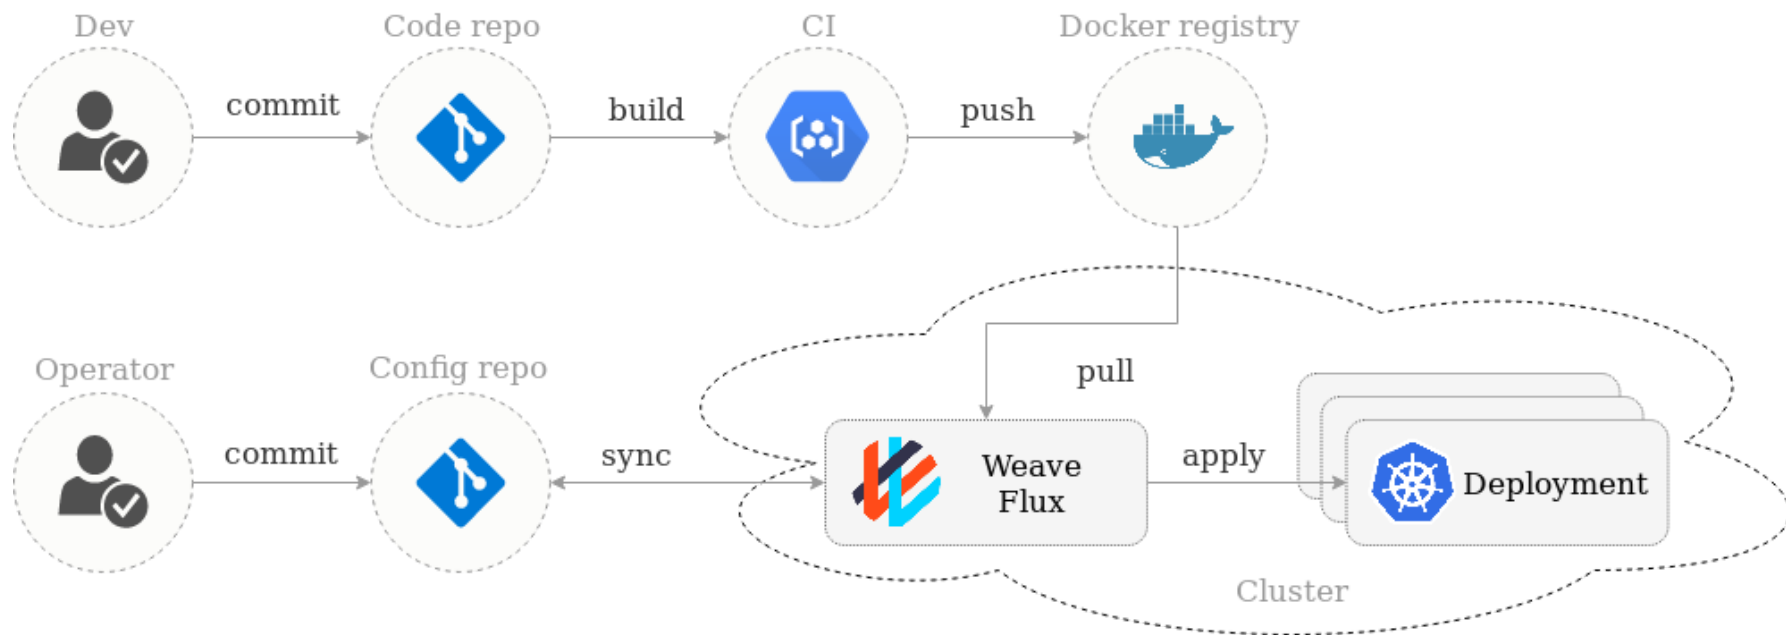
- Git is one of the most common tool for source code management
- GitOps is an operational and pipeline pattern that establishes the same concept not just for source code but for every operational aspect of your system
- All system and environments configuration is defined through Git, and the process of applying changes to the infrastructure follows the same practices and quality control procedures that a change in software does. Making it possible to review, check and promote across branches/environments any change that is made to a system.



# IaC tools vs GitOps

- IaC originated the concept of keeping infrastructure config versioned, backed up and reproducible from source code
- Kubernetes is almost completely declarative, combined with the immutable container, it is possible to extend the concepts to manage applications and their operating system.
- GitOps philosophy is to test, deploy, rollback, rollforward with a complete audit trail all from Git
- IaC to provision servers, GitOps to manage Kubernetes

# Weaveworks Flux - The GitOps Kubernetes operator



# Options

- **Argoproj** Makes use of CRD to define pipelines. Cloud Native. GitOps.
- **Jenkin X** GitOps and runs on Kubernetes with a very powerfull CLI **jx**
- **InfraBox** CIOps that runs on Kubernetes with Docker jobs
- **Jenkins** CIOps with Kubernetes plugins
- **GitLab - Auto DevOps** Built-in container registry and Kubernetes support
- **Spinnaker** Multi-cloud continuos delivery platform
- **Skaffold** Tool for local and remote developing, but can be used in CI/CD.

# Getting started with GitOps

Fork & clone the repository

```
git clone https://github.com/stacc-as/gitops-workshop
```



# 1. Install the flux operator on your cluster

- Install flux

```
cd gitops-tutorial  
kubectl apply -f ./flux/
```

- Alongside flux, we have also installed a memcache service for flux to use, and appropriate permissions for the flux operator to act on the cluster itself through a role and rolebinding

```
kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
flux-7cddb59dd4-kftxn	1/1	Running	0	15d
flux-helm-operator-7f8467486d-hx74w	1/1	Running	0	22d

## 2. Sharing the operator's public key with the remote git repository

- Download the fluxctl CLI

```
wget https://github.com/weaveworks/flux/releases/download/1.11.0/fluxctl_linux_...  
/usr/local/bin/fluxctl \  
&& chmod +x /usr/local/bin/fluxctl
```

- And retrieve the public key from the operator

```
fluxctl identity  
ssh-rsa AAAAB3NzaC1yc2EAAAADAQABAAQBAQC105 [...]
```

- Or via kubectl

```
kubectl logs deployment/flux | grep identity.pub | cut -d '"' -f2
```

- Add the key to the remote repository

# 3. Configure the operator to read from your repository

- Edit the `flux/flux-deployment.yaml` and update the following values

```
# Configure git repository
--git-url=git@github.com:stacc-as/gitops-workshop.git
# Configure git branch
--git-branch=master
# Configure the path for flux to read
--git-path=deploy/kubernetes
# Configure the git user
--git-user=gitops-workshop
# Configure the git email
--git-email=flux@stacc.com
# Configure git sync label, must end with '-sync'
--git-label=flux-sync
```

- Reapply the configuration to the cluster

```
kubectl apply -f flux/deployment.yaml
```

## 4. Add a yaml file

The operator is set up to react to changes in the `deploy/kubernetes` folder, we need to provide it with some configuration for it to synchronise

```
● ● ●  
cp examples/podinfo-dep.yaml deploy/kubernetes  
git add .  
git commit -m "Add podinfo to cluster"  
git push
```



# 5. Watch for the magic

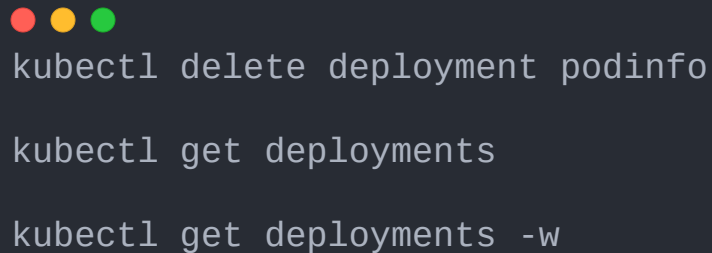
The flux operator synchronize every 5 minutes, the `podinfo` deployment should be up and running after a little while



```
kubectl get pods -w
```

# 6. Lets fuck things up!

Wait for the flux to sync, it will detect the missing deployment and recreate it from the last good state



```
kubectl delete deployment podinfo  
  
kubectl get deployments  
  
kubectl get deployments -w
```

# 7. Modify the configuration

```

$ kubectl get deployment podinfo

NAME          DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
podinfo       1         1         1            1           11m

```

- Change the replicas from 1 to 4 in `deploy/kubernetes-podinfo-dep.yaml`
- Commit the changes
- Wait for flux to sync and watch changes take action

```

kubectl get deployment podinfo

```

# 8. Automate deployment



```
annotations:  
  flux.weave.works/tag.podinfo: glob:*  
  flux.weave.works/automated: 'true'
```

# Questions?

